

MLP & CNN for Google Street View House Numbers

January 3, 2021

1 Capstone Project

1.1 Image classifier for the Street View House Numbers dataset

The task is to develop MLP and CNN models in order to verify their performance as well as generalization on the test set. The assumption is that CNN will have less parameters requiring a bit longer training time, but resulting in better generalization on unseen set.

Both models' types will use the same optimizer and loss, as well as metrics.

```
[1]: import numpy as np
import tensorflow as tf
import matplotlib.pyplot as plt

from scipy.io import loadmat
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, BatchNormalization, Dense, Dropout,
    ↳ Flatten, MaxPooling2D
from tensorflow.keras.regularizers import l1_l2
from tensorflow.keras.activations import relu, softmax
from tensorflow.keras.initializers import he_uniform, ones
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import SparseCategoricalCrossentropy
from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping,
    ↳ ReduceLROnPlateau
```

1.1.1 Load and inspect the sets

```
[3]: train_data = loadmat('./project_svhn_classification/assets/train_32x32.mat')
test_data = loadmat('./project_svhn_classification/assets/test_32x32.mat')
```

```
[4]: X_train, y_train = train_data['X'], train_data['y']
X_test, y_test = test_data['X'], test_data['y']

print('X training shape: ', X_train.shape)
print('y train shape: ', y_train.shape)
print('\nUniq labels in the train set: ', np.unique(y_train))
print('Uniq labels in the test set: ', np.unique(y_test))
```

```
X training shape: (32, 32, 3, 73257)
y train shape: (73257, 1)
```

```
Uniq labels in the train set: [ 1  2  3  4  5  6  7  8  9 10]
Uniq labels in the test set: [ 1  2  3  4  5  6  7  8  9 10]
```

Replace the label “10” with “0”

```
[5]: y_train[y_train == 10] = 0
      y_test[y_test == 10] = 0

      print('\nUniq labels in the train set: ', np.unique(y_train))
      print('Uniq labels in the test set: ', np.unique(y_test))
```

```
Uniq labels in the train set: [0 1 2 3 4 5 6 7 8 9]
Uniq labels in the test set: [0 1 2 3 4 5 6 7 8 9]
```

Move the samples axis as first

```
[10]: X_train = np.moveaxis(X_train, -1, 0)
      X_test = np.moveaxis(X_test, -1, 0)

      print('X training shape: ', X_train.shape)
      print('X testing shape: ', X_test.shape)
```

```
X training shape: (73257, 32, 32, 3)
X testing shape: (26032, 32, 32, 3)
```

Display images from the “train” set

```
[21]: def display_sample_images(data_set, labels_set):
      fig, ax = plt.subplots(nrows=1, ncols=10, figsize=(20, 10))

      for i in range(10):
          idx = np.random.choice(range(data_set.shape[0]))
          ax[i].set_axis_off()
          ax[i].imshow(data_set[idx], interpolation='nearest')
          ax[i].set_title(f'Label: {labels_set[idx][0]}')

      plt.subplots_adjust(wspace=1)
```

```
[22]: display_sample_images(X_train, y_train)
```



Display images from the “test” set

```
[24]: display_sample_images(X_test, y_test)
```



Normalize data


```
[25]: X_train_inputs = np.mean(X_train, axis=3, keepdims=True) / 255.  
X_test_inputs = np.mean(X_test, axis=3, keepdims=True) / 255.  
  
print('X training inputs shape: ', X_train_inputs.shape)  
print('X testing inputs shape: ', X_test_inputs.shape)
```

X training inputs shape: (73257, 32, 32, 1)

X testing inputs shape: (26032, 32, 32, 1)

Display normalized images - train set

```
[28]: display_sample_images(X_train_inputs, y_train)
```



1.1.2 The MLP model

```
[31]: def build_mlp_model(input_shape):  
    model = Sequential([  
        Flatten(input_shape=input_shape, name='flatten'),  
        Dense(units=128,  
              activation=relu,  
              kernel_initializer=he_uniform(),  
              bias_initializer=he_uniform(),  
              name='dense_1'),  
        Dense(units=64,  
              activation=relu,  
              kernel_initializer=he_uniform(),  
              bias_initializer=he_uniform(),  
              name='dense_2'),  
        Dense(units=32,  
              activation=relu,  
              kernel_initializer=he_uniform(),  
              bias_initializer=he_uniform(),
```

```

        name='dense_3'),
Dense(units=10,
      activation=softmax,
      kernel_initializer=he_uniform(),
      bias_initializer=he_uniform(),
      name='dense_output')
])
model.compile(
    optimizer=Adam(),
    loss=SparseCategoricalCrossentropy(),
    metrics=['accuracy'])
return model

```

```

[61]: mlp_model = build_mlp_model(input_shape=X_train_inputs[0].shape)
      mlp_model.summary()

```

Model: "sequential_7"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 1024)	0
dense_1 (Dense)	(None, 128)	131200
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 32)	2080
dense_output (Dense)	(None, 10)	330

=====
 Total params: 141,866
 Trainable params: 141,866
 Non-trainable params: 0
 =====

Learning callbacks

```

[33]: best_mlp_model_cb = ModelCheckpoint(filepath='most_accurate_model/
      ↪model_{epoch}',

      monitor='val_accuracy',
      save_best_only=True,
      save_weights_only=True,
      mode='max',
      verbose=1)

early_stopping_cb = EarlyStopping(monitor='val_loss',
      min_delta=1e-2,
      patience=4,

```

```

        verbose=1)

lr_reduction_cb = ReduceLROnPlateau(monitor='val_loss',
                                     factor=1e-4,
                                     patience=2,
                                     mode='min',
                                     verbose=1)

```

Train MLP model

```

[62]: mlp_history = mlp_model.fit(x=X_train_inputs,
                                y=y_train,
                                epochs=50,
                                batch_size=64,
                                validation_data=(X_test_inputs, y_test),
                                callbacks=[
                                    best_mlp_model_cb,
                                    early_stopping_cb,
                                    lr_reduction_cb,
                                ])

```

Epoch 1/50

1145/1145 [=====] - 2s 2ms/step - loss: 2.2490 - accuracy: 0.1855 - val_loss: 1.9035 - val_accuracy: 0.3633

Epoch 00001: val_accuracy did not improve from 0.74977

Epoch 2/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.6686 - accuracy: 0.4295 - val_loss: 1.4014 - val_accuracy: 0.5515

Epoch 00002: val_accuracy did not improve from 0.74977

Epoch 3/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.3076 - accuracy: 0.5754 - val_loss: 1.3510 - val_accuracy: 0.5759

Epoch 00003: val_accuracy did not improve from 0.74977

Epoch 4/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.2059 - accuracy: 0.6158 - val_loss: 1.2785 - val_accuracy: 0.6007

Epoch 00004: val_accuracy did not improve from 0.74977

Epoch 5/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.1502 - accuracy: 0.6370 - val_loss: 1.2570 - val_accuracy: 0.6099

Epoch 00005: val_accuracy did not improve from 0.74977

Epoch 6/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.1224 -

accuracy: 0.6480 - val_loss: 1.2407 - val_accuracy: 0.6117

Epoch 00006: val_accuracy did not improve from 0.74977

Epoch 7/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.0867 -
accuracy: 0.6595 - val_loss: 1.2329 - val_accuracy: 0.6280

Epoch 00007: val_accuracy did not improve from 0.74977

Epoch 8/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.0533 -
accuracy: 0.6739 - val_loss: 1.1785 - val_accuracy: 0.6425

Epoch 00008: val_accuracy did not improve from 0.74977

Epoch 9/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.0299 -
accuracy: 0.6795 - val_loss: 1.1726 - val_accuracy: 0.6483

Epoch 00009: val_accuracy did not improve from 0.74977

Epoch 10/50

1145/1145 [=====] - 2s 2ms/step - loss: 1.0145 -
accuracy: 0.6869 - val_loss: 1.1777 - val_accuracy: 0.6423

Epoch 00010: val_accuracy did not improve from 0.74977

Epoch 11/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9961 -
accuracy: 0.6945 - val_loss: 1.1182 - val_accuracy: 0.6626

Epoch 00011: val_accuracy did not improve from 0.74977

Epoch 12/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9920 -
accuracy: 0.6932 - val_loss: 1.1509 - val_accuracy: 0.6524

Epoch 00012: val_accuracy did not improve from 0.74977

Epoch 13/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9865 -
accuracy: 0.6959 - val_loss: 1.0840 - val_accuracy: 0.6718

Epoch 00013: val_accuracy did not improve from 0.74977

Epoch 14/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9814 -
accuracy: 0.6964 - val_loss: 1.1432 - val_accuracy: 0.6481

Epoch 00014: val_accuracy did not improve from 0.74977

Epoch 15/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9645 -
accuracy: 0.7019 - val_loss: 1.0925 - val_accuracy: 0.6694

Epoch 00015: val_accuracy did not improve from 0.74977

Epoch 00015: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-07.

Epoch 16/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9480 - accuracy: 0.7076 - val_loss: 1.0825 - val_accuracy: 0.6728

Epoch 00016: val_accuracy did not improve from 0.74977

Epoch 17/50

1145/1145 [=====] - 2s 2ms/step - loss: 0.9340 - accuracy: 0.7139 - val_loss: 1.0755 - val_accuracy: 0.6753

Epoch 00017: val_accuracy did not improve from 0.74977

Epoch 00017: early stopping

```
[36]: def evaluate_model(model, name, set_inputs, labels):
      test_loss, test_accuracy = model.evaluate(x=set_inputs, y=labels, verbose=0)
      print(f'Evaluating model {name}')
      print(f'* Test accuracy: {np.round(test_accuracy, 3)}')
      print(f'* Test loss: {np.round(test_loss, 3)}')
```

```
[63]: evaluate_model(mlp_model, name='MLP', set_inputs=X_test_inputs, labels=y_test)
```

Evaluating model MLP

* Test accuracy: 0.675


* Test loss: 1.076

MLP model's performance: accuracy and sparse categorical crossentropy loss over epochs

```
[59]: def plot_history_performance(history):
      fig = plt.figure(figsize=(12, 5))
      fig.add_subplot(121)
      plt.plot(history['accuracy'])
      plt.plot(history['val_accuracy'])
      plt.title('Accuracy over Epochs')
      plt.xlabel('Epoch')
      plt.ylabel('Accuracy')
      plt.legend(['Training', 'Validation'], loc='lower right')

      fig.add_subplot(122)
      plt.plot(history['loss'])
      plt.plot(history['val_loss'])
      plt.title('Loss over Epochs')
      plt.xlabel('Epoch')
      plt.ylabel('Loss')
      plt.legend(['Training', 'Validation'], loc='upper right')
```

```
[64]: plot_history_performance(mlp_history.history)
```



1.1.3 The CNN model

```
[44]: def build_cnn_model(input_shape):
    model = Sequential([
        Conv2D(filters=16,
               kernel_size=(3, 3),
               input_shape=input_shape,
               padding='same',
               activation=relu,
               kernel_initializer=he_uniform(),
               kernel_regularizer=l1_l2(l1=1e-3, l2=1e-3),
               bias_initializer=he_uniform(),
               bias_regularizer=l1_l2(l1=1e-3, l2=1e-3),
               name='conv2d_1'),
        Conv2D(filters=8,
               kernel_size=(3, 3),
               padding='same',
               activation=relu,
               kernel_initializer=he_uniform(),
               kernel_regularizer=l1_l2(l1=1e-3, l2=1e-3),
               bias_initializer=he_uniform(),
               bias_regularizer=l1_l2(l1=1e-3, l2=1e-3),
               name='conv2d_2'),
        MaxPooling2D(pool_size=(3, 3), name='max_pooling_2d_1'),
        BatchNormalization(name='batch_normalization_1'),
        Dropout(rate=0.3, name='dropout_1'),
        Conv2D(filters=8,
               kernel_size=(3, 3),
               padding='same',
```



```

        activation=relu,
        kernel_initializer=he_uniform(),
        kernel_regularizer=l1_l2(l1=1e-3, l2=1e-3),
        bias_initializer=he_uniform(),
        bias_regularizer=l1_l2(l1=1e-3, l2=1e-3),
        name='conv2d_3'),
MaxPooling2D(pool_size=(3, 3), name='max_pooling_2d_2'),
Flatten(name='flatten'),
Dense(units=32,
      activation=relu,
      kernel_initializer=he_uniform(),
      kernel_regularizer=l1_l2(l1=1e-3, l2=1e-3),
      bias_initializer=he_uniform(),
      bias_regularizer=l1_l2(l1=1e-3, l2=1e-3),
      name='dense_1'),
Dropout(0.3, name='dropout_2'),
Dense(units=16,
      activation=relu,
      kernel_initializer=he_uniform(),
      bias_initializer=he_uniform(),
      kernel_regularizer=l1_l2(l1=1e-3, l2=1e-3),
      name='dense_2'),
Dense(units=10,
      activation=softmax,
      kernel_initializer=he_uniform(),
      kernel_regularizer=l1_l2(l1=1e-3, l2=1e-3),
      bias_initializer=he_uniform(),
      bias_regularizer=l1_l2(l1=1e-3, l2=1e-3),
      name='dense_output')
])
model.compile(
    optimizer=Adam(),
    loss=SparseCategoricalCrossentropy(),
    metrics=['accuracy'])
return model

```

```

[45]: cnn_model = build_cnn_model(input_shape=X_train_inputs[0].shape)
      cnn_model.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 16)	160
conv2d_2 (Conv2D)	(None, 32, 32, 8)	1160

max_pooling_2d_1 (MaxPooling)	(None, 10, 10, 8)	0

batch_normalization_1 (Batch Normalization)	(None, 10, 10, 8)	32

dropout_1 (Dropout)	(None, 10, 10, 8)	0

conv2d_3 (Conv2D)	(None, 10, 10, 8)	584

max_pooling_2d_2 (MaxPooling)	(None, 3, 3, 8)	0

flatten (Flatten)	(None, 72)	0

dense_1 (Dense)	(None, 32)	2336

dropout_2 (Dropout)	(None, 32)	0

dense_2 (Dense)	(None, 16)	528

dense_output (Dense)	(None, 10)	170
=====		
Total params: 4,970		
Trainable params: 4,954		
Non-trainable params: 16		

```
[46]: best_cnn_model_cb = ModelCheckpoint(filepath='most_accurate_cnn_model/
      ↪model_{epoch}',
      monitor='val_accuracy',
      save_best_only=True,
      save_weights_only=True,
      mode='max',
      verbose=1)
```

```
[49]: cnn_history = cnn_model.fit(x=X_train_inputs,
      y=y_train,
      epochs=50,
      batch_size=64,
      validation_data=(X_test_inputs, y_test),
      callbacks=[
          best_cnn_model_cb,
          early_stopping_cb,
          lr_reduction_cb,
      ])
```

Epoch 1/50
 1145/1145 [=====] - 32s 27ms/step - loss: 2.6862 -
 accuracy: 0.1910 - val_loss: 2.2098 - val_accuracy: 0.2995

Epoch 00001: val_accuracy did not improve from 0.75192
Epoch 2/50
1145/1145 [=====] - 32s 28ms/step - loss: 1.9141 - accuracy: 0.4096 - val_loss: 1.5982 - val_accuracy: 0.5423

Epoch 00002: val_accuracy did not improve from 0.75192
Epoch 3/50
1145/1145 [=====] - 32s 28ms/step - loss: 1.6568 - accuracy: 0.5145 - val_loss: 1.4570 - val_accuracy: 0.6044

Epoch 00003: val_accuracy did not improve from 0.75192
Epoch 4/50
1145/1145 [=====] - 32s 28ms/step - loss: 1.4784 - accuracy: 0.5929 - val_loss: 1.2287 - val_accuracy: 0.6999

Epoch 00004: val_accuracy did not improve from 0.75192
Epoch 5/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.3742 - accuracy: 0.6366 - val_loss: 1.1378 - val_accuracy: 0.7372

Epoch 00005: val_accuracy did not improve from 0.75192
Epoch 6/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.3066 - accuracy: 0.6630 - val_loss: 1.0928 - val_accuracy: 0.7495

Epoch 00006: val_accuracy did not improve from 0.75192
Epoch 7/50
1145/1145 [=====] - 34s 29ms/step - loss: 1.2490 - accuracy: 0.6876 - val_loss: 1.0885 - val_accuracy: 0.7523

Epoch 00007: val_accuracy improved from 0.75192 to 0.75234, saving model to most_accurate_cnn_model/model_7
Epoch 8/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.2205 - accuracy: 0.6988 - val_loss: 0.9894 - val_accuracy: 0.7888

Epoch 00008: val_accuracy improved from 0.75234 to 0.78880, saving model to most_accurate_cnn_model/model_8
Epoch 9/50
1145/1145 [=====] - 34s 29ms/step - loss: 1.1988 - accuracy: 0.7052 - val_loss: 0.9674 - val_accuracy: 0.7985

Epoch 00009: val_accuracy improved from 0.78880 to 0.79848, saving model to most_accurate_cnn_model/model_9
Epoch 10/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.1810 - accuracy: 0.7141 - val_loss: 0.9648 - val_accuracy: 0.7917

Epoch 00010: val_accuracy did not improve from 0.79848
Epoch 11/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.1691 - accuracy: 0.7143 - val_loss: 0.9492 - val_accuracy: 0.8005

Epoch 00011: val_accuracy improved from 0.79848 to 0.80051, saving model to most_accurate_cnn_model/model_11
Epoch 12/50
1145/1145 [=====] - 32s 28ms/step - loss: 1.1568 - accuracy: 0.7212 - val_loss: 0.9296 - val_accuracy: 0.8040

Epoch 00012: val_accuracy improved from 0.80051 to 0.80405, saving model to most_accurate_cnn_model/model_12
Epoch 13/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.1477 - accuracy: 0.7225 - val_loss: 0.9331 - val_accuracy: 0.8002

Epoch 00013: val_accuracy did not improve from 0.80405
Epoch 14/50
1145/1145 [=====] - 35s 30ms/step - loss: 1.1409 - accuracy: 0.7243 - val_loss: 0.9837 - val_accuracy: 0.8022

Epoch 00014: val_accuracy did not improve from 0.80405

Epoch 00014: ReduceLROnPlateau reducing learning rate to 1.0000000474974514e-07.
Epoch 15/50
1145/1145 [=====] - 36s 31ms/step - loss: 1.1222 - accuracy: 0.7327 - val_loss: 0.9113 - val_accuracy: 0.8137

Epoch 00015: val_accuracy improved from 0.80405 to 0.81369, saving model to most_accurate_cnn_model/model_15
Epoch 16/50
1145/1145 [=====] - 37s 32ms/step - loss: 1.1228 - accuracy: 0.7334 - val_loss: 0.9105 - val_accuracy: 0.8139

Epoch 00016: val_accuracy improved from 0.81369 to 0.81388, saving model to most_accurate_cnn_model/model_16
Epoch 17/50
1145/1145 [=====] - 35s 31ms/step - loss: 1.1197 - accuracy: 0.7327 - val_loss: 0.9096 - val_accuracy: 0.8139

Epoch 00017: val_accuracy did not improve from 0.81388
Epoch 18/50
1145/1145 [=====] - 33s 29ms/step - loss: 1.1208 - accuracy: 0.7321 - val_loss: 0.9085 - val_accuracy: 0.8144

Epoch 00018: val_accuracy improved from 0.81388 to 0.81438, saving model to most_accurate_cnn_model/model_18

Epoch 19/50
 1145/1145 [=====] - 36s 31ms/step - loss: 1.1202 -
 accuracy: 0.7323 - val_loss: 0.9080 - val_accuracy: 0.8144


Epoch 00019: val_accuracy did not improve from 0.81438
 Epoch 00019: early stopping

```
[50]: evaluate_model(cnn_model, name='CNN', set_inputs=X_test_inputs, labels=y_test)
```

Evaluating model CNN
 * Test accuracy: 0.814
 * Test loss: 0.908

MLP model's performance: accuracy and sparse categorical crossentropy loss over epochs

```
[65]: plot_history_performance(cnn_history.history)
```



1.1.4 Load best models

```
[55]: best_mlp_model = build_mlp_model(input_shape=X_train_inputs[0].shape)
best_mlp_model.load_weights(tf.train.latest_checkpoint('most_accurate_model'))
evaluate_model(best_mlp_model, name='MLP', set_inputs=X_test_inputs,
↳ labels=y_test)
```

Evaluating model MLP
 * Test accuracy: 0.75
 * Test loss: 0.857

```
[56]: best_cnn_model = build_cnn_model(input_shape=X_train_inputs[0].shape)
best_cnn_model.load_weights(tf.train.
    ↳latest_checkpoint('most_accurate_cnn_model'))
evaluate_model(best_cnn_model, name='CNN', set_inputs=X_test_inputs,
    ↳labels=y_test)
```

```
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.iter
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_1
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_2
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.decay
WARNING:tensorflow:Unresolved object in checkpoint:
(root).optimizer.learning_rate
WARNING:tensorflow:A checkpoint was restored (e.g. tf.train.Checkpoint.restore
or tf.keras.Model.load_weights) but not all checkpointed values were used. See
above for specific issues. Use expect_partial() on the load status object, e.g.
tf.train.Checkpoint.restore(...).expect_partial(), to silence these warnings, or
use assert_consumed() to make the check explicit. See
https://www.tensorflow.org/guide/checkpoint#loading\_mechanics for details.
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.iter
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_1
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_2
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.decay
WARNING:tensorflow:Unresolved object in checkpoint:
(root).optimizer.learning_rate
WARNING:tensorflow:A checkpoint was restored (e.g. tf.train.Checkpoint.restore
or tf.keras.Model.load_weights) but not all checkpointed values were used. See
above for specific issues. Use expect_partial() on the load status object, e.g.
tf.train.Checkpoint.restore(...).expect_partial(), to silence these warnings, or
use assert_consumed() to make the check explicit. See
https://www.tensorflow.org/guide/checkpoint#loading\_mechanics for details.
Evaluating model CNN
* Test accuracy: 0.814
* Test loss: 0.908
```

Compare predictions using the test set

```
[75]: fig, ax = plt.subplots(nrows=15, ncols=3, figsize=(18, 50))

x_values = list(range(10))

for i in range(15):
    idx = np.random.choice(range(X_test_inputs.shape[0]))
    ax[i, 0].set_axis_off()
    ax[i, 0].imshow(X_test[idx], interpolation='nearest')
    ax[i, 0].set_title(f'#{idx} Class: {y_test[idx][0]}')

    img_input = X_test_inputs[idx][np.newaxis, ...]
```

```
mlp_predictions = best_mlp_model.predict(img_input)
mlp_prediction = np.argmax(mlp_predictions)
ax[i, 1].set_title(f'MLP Pred: {mlp_prediction}')
ax[i, 1].bar(x=x_values, height=mlp_predictions[0], tick_label=x_values)

cnn_predictions = best_cnn_model.predict(img_input)
cnn_prediction = np.argmax(cnn_predictions)
ax[i, 2].set_title(f'CNN Pred: {cnn_prediction}')
ax[i, 2].bar(x=x_values, height=cnn_predictions[0], tick_label=x_values)

plt.subplots_adjust(hspace=.5)
```

